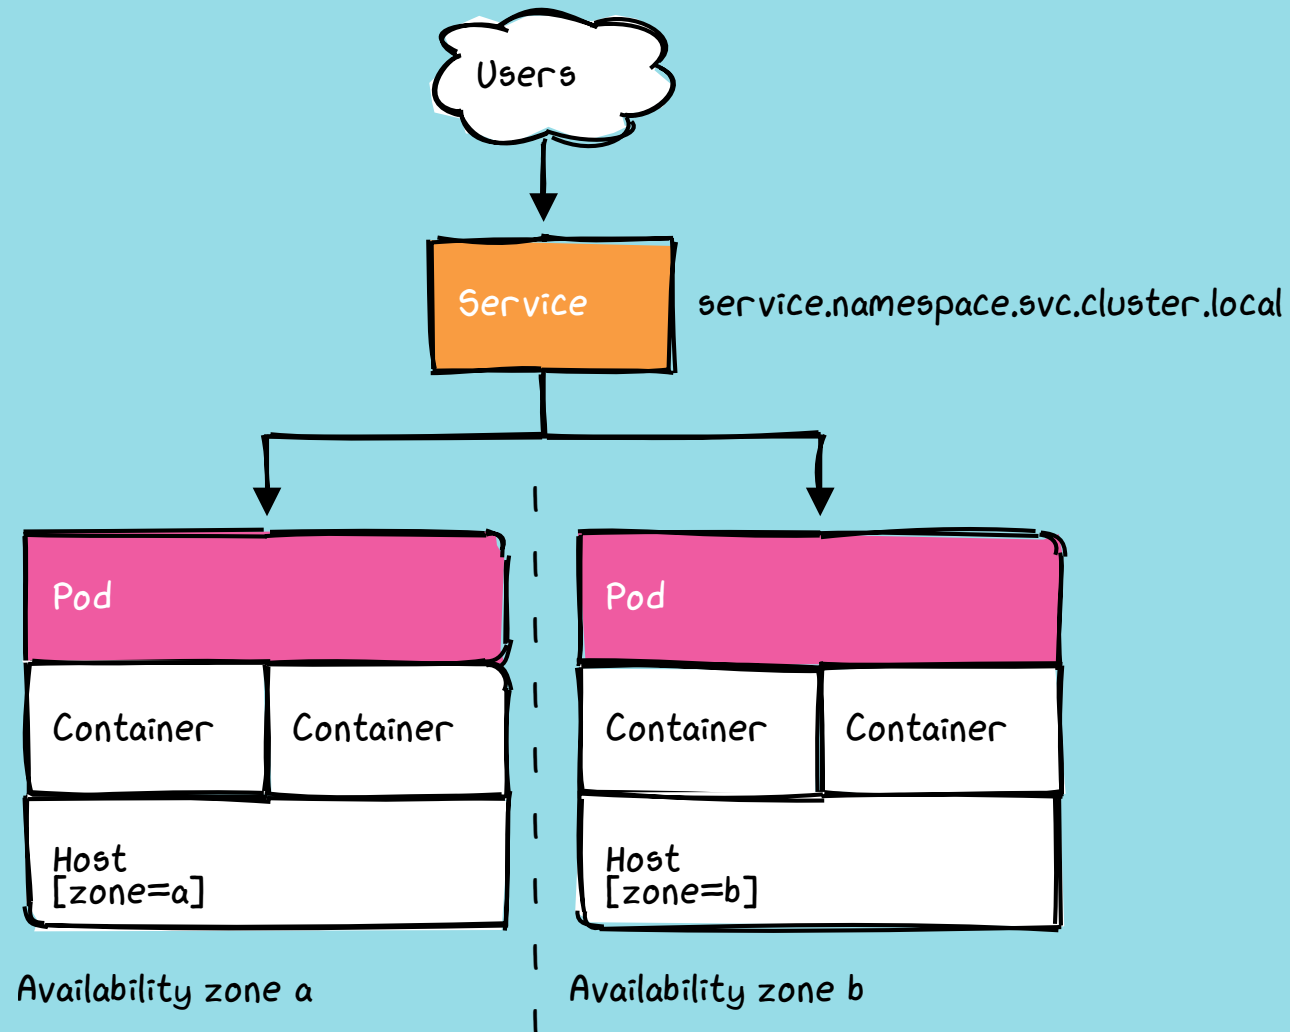


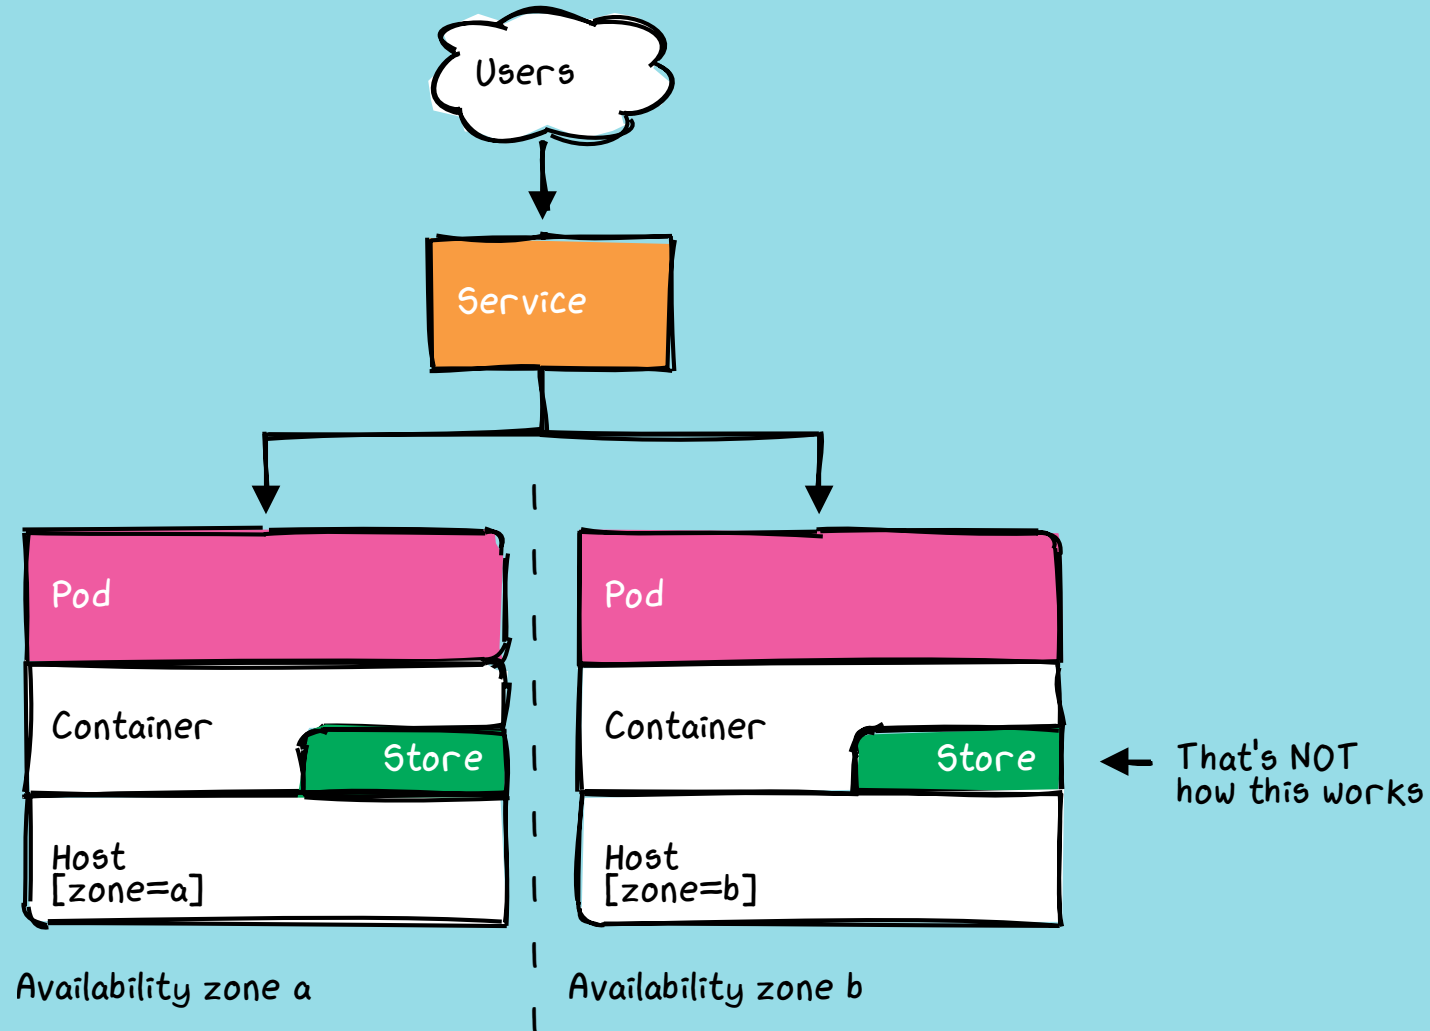
STATEFUL KUBERNETES

PAST, PRESENT, FUTURE

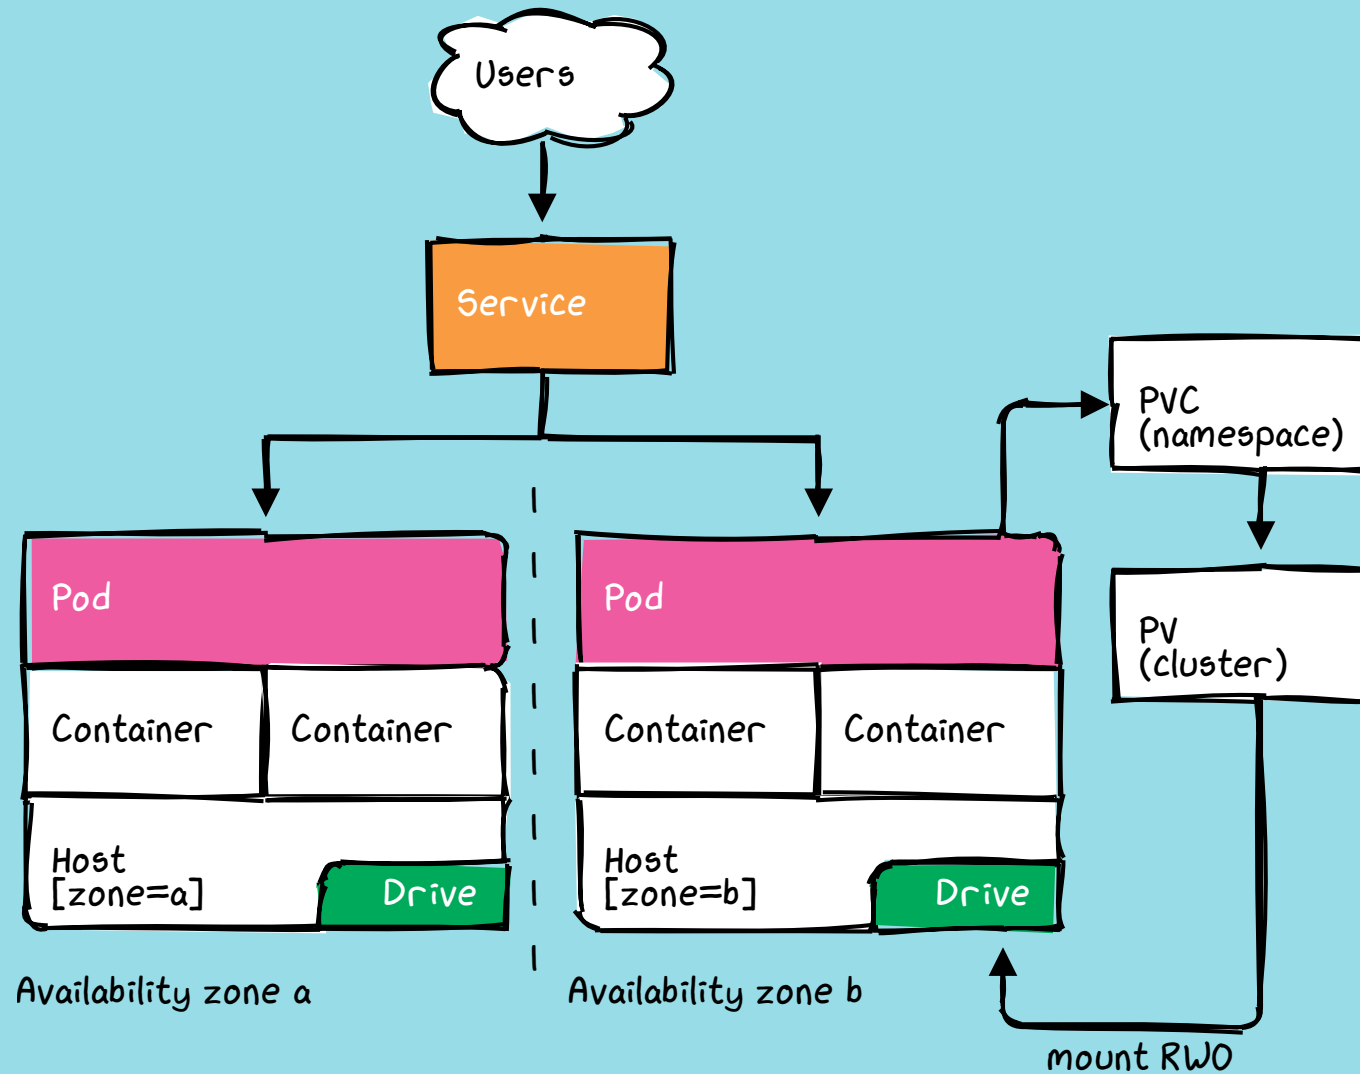
STATELESS KUBERNETES



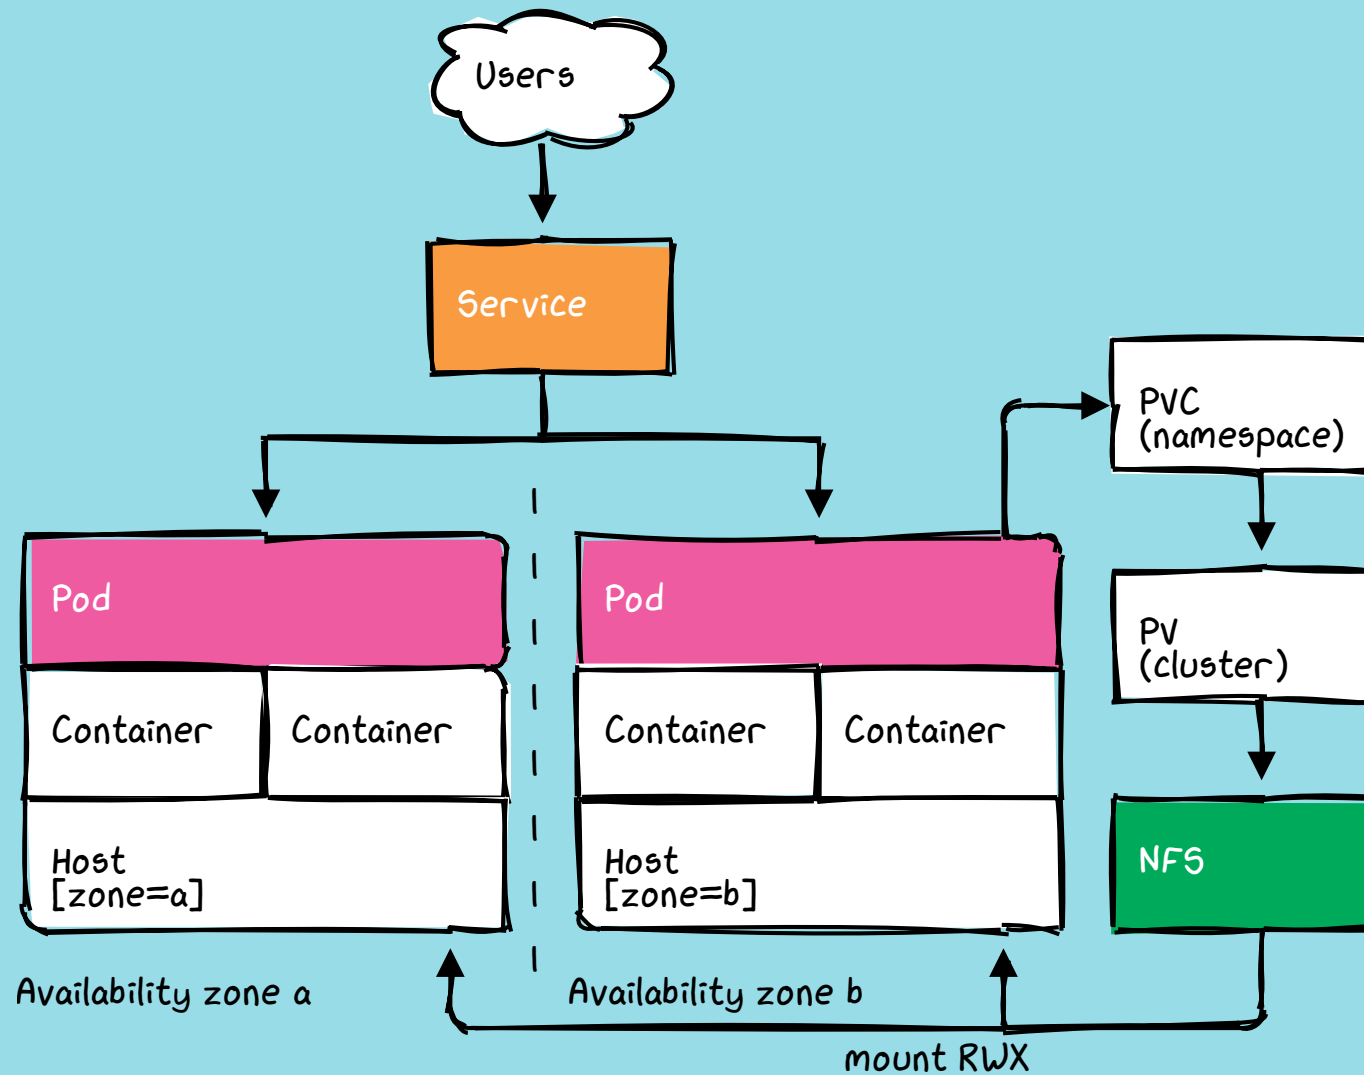
CONTAINER STORAGE FABLE



CONTAINER BLOCK STORAGE



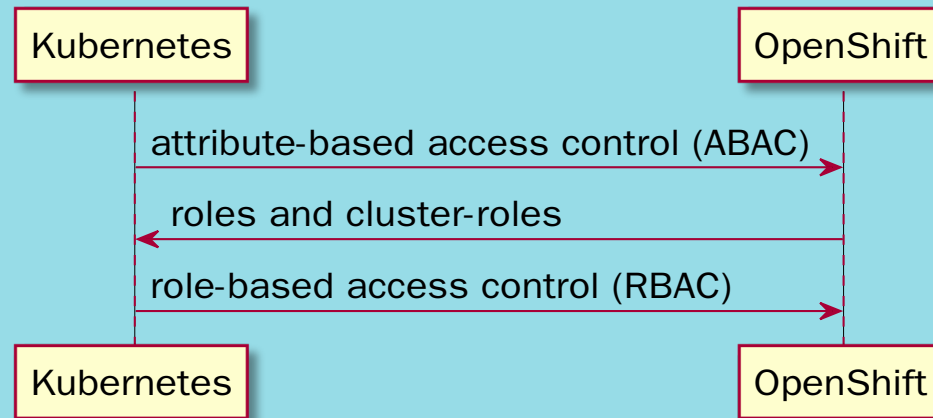
CONTAINER NETWORK FILE STORAGE




TIMELINES

IDENTITY, NETWORK, STORAGE AND STATE

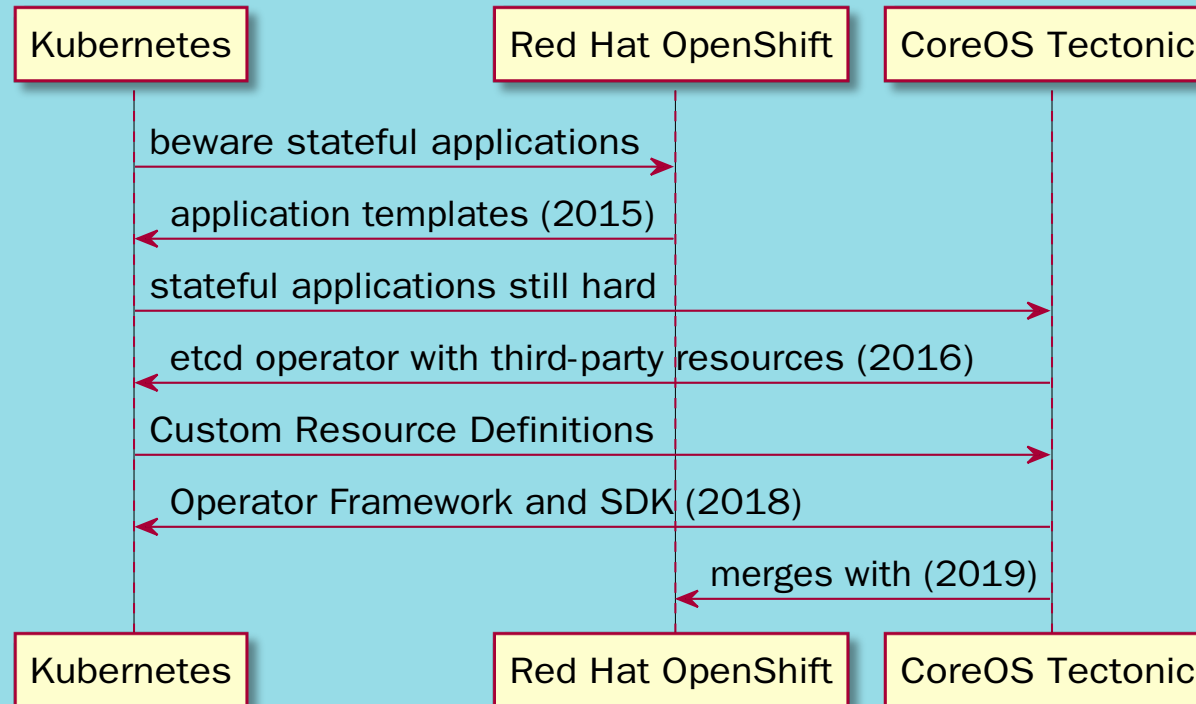
IDENTITY




NETWORK



STORAGE AND STATE

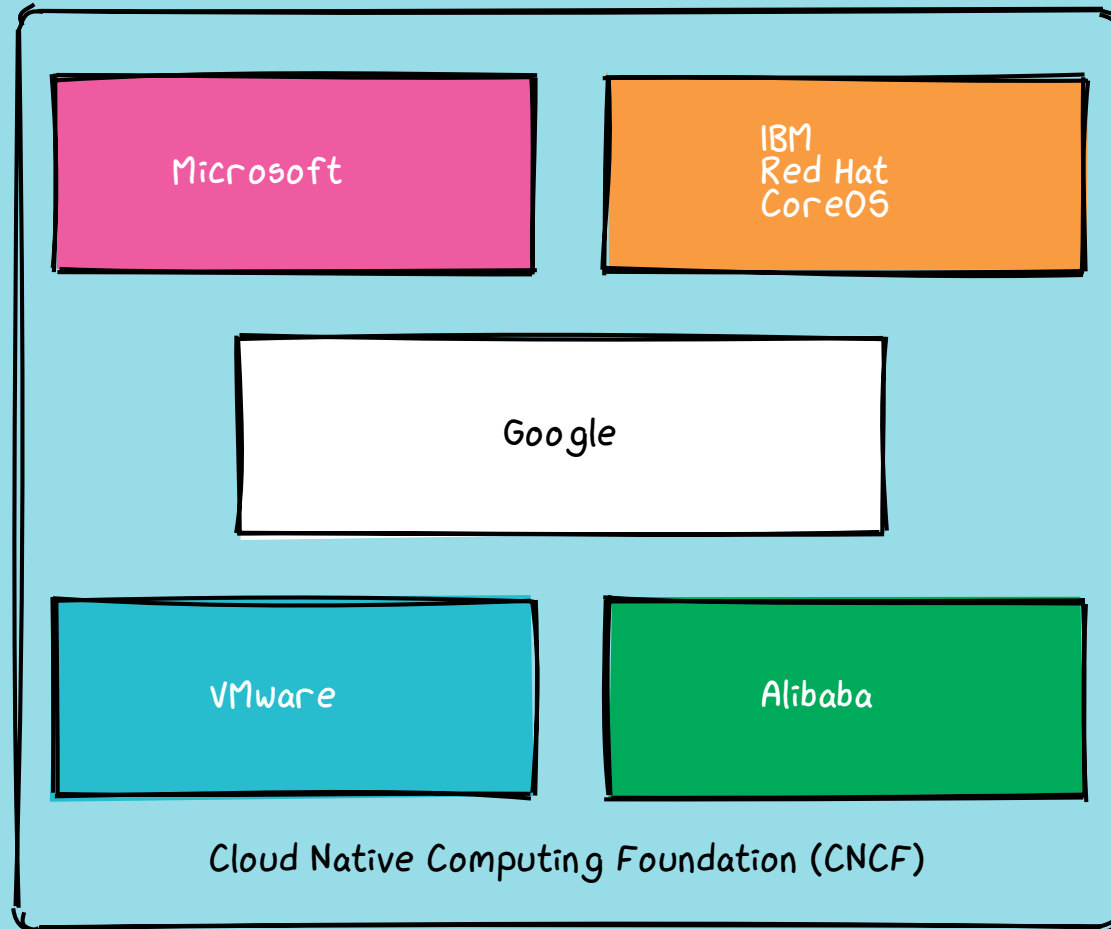


REBUILDING THE DATACENTRE



ECOSYSTEM RESILIENCE

SHARED OWNERSHIP



PRECONDITIONS

open source

neutral IP ownership

extensibility

PRECONDITIONS

open source – Apache
neutral IP ownership
extensibility

A suitably permissive license is a necessary but not sufficient precondition.

PRECONDITIONS

open source – Apache

neutral IP ownership – CNCF

extensibility


PRECONDITIONS

K8s would **not** be what it is today without neutral IP ownership... K8s has spawned an entire ecosystem **because** it can be used by consuming projects/products without fear.

– Matt Klein

Source: @mattklein123 on Twitter, 17 February 2020.

PRECONDITIONS

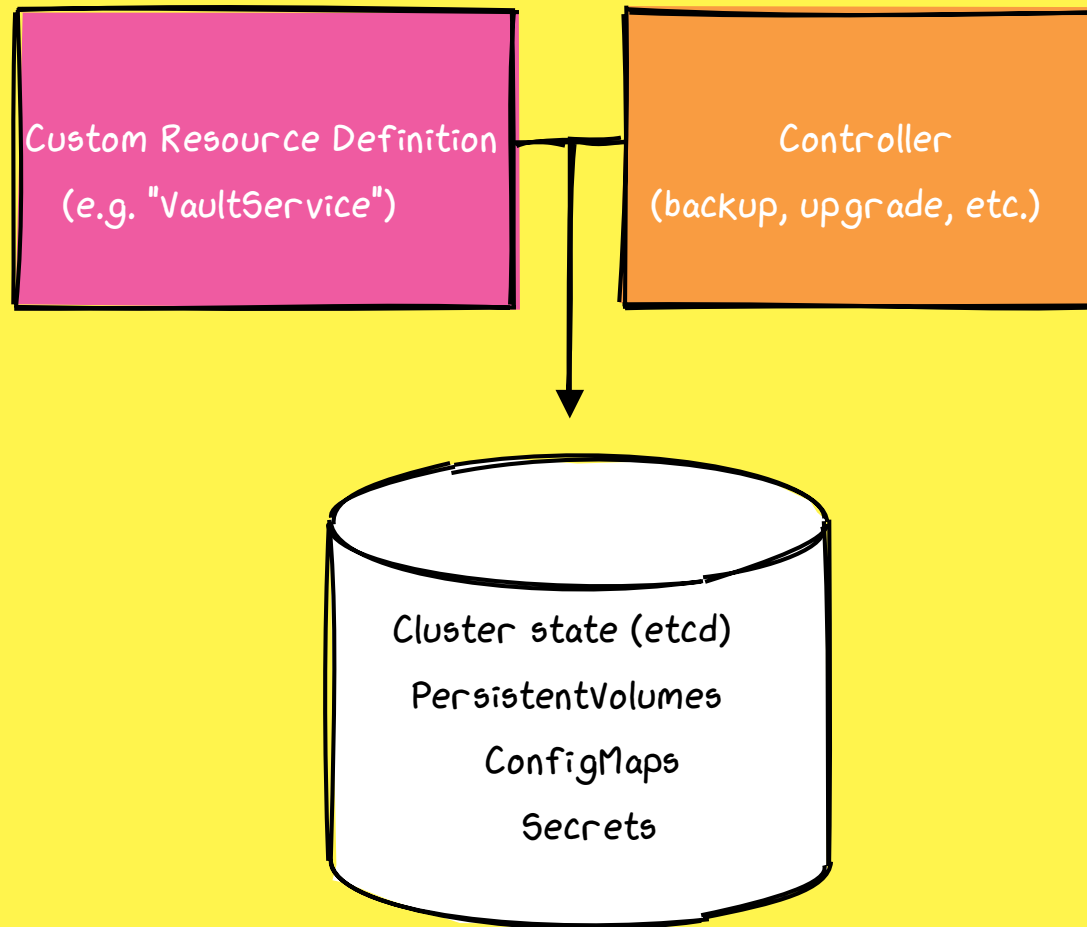


Controllers add custom processing to the core reconciliation loop.

When paired with custom resource definitions, they are known as **operators**.

THE OPERATOR LIFE

OPERATORS



CORPORATE SPONSORS

Vault operator



MySQL operator



PostgreSQL operator



STATEFUL WORKLOADS

A pink speech bubble with a black outline, pointing downwards on the left side. It contains the text "Stateless is Easy, Stateful is Hard." and "- Brandon Philips (2016)".


Stateless is Easy, Stateful is Hard.

– Brandon Philips (2016)

Source: coreos.com/blog/introducing-operators.html

SOUL-SEARCHING AT THE BAZAAR

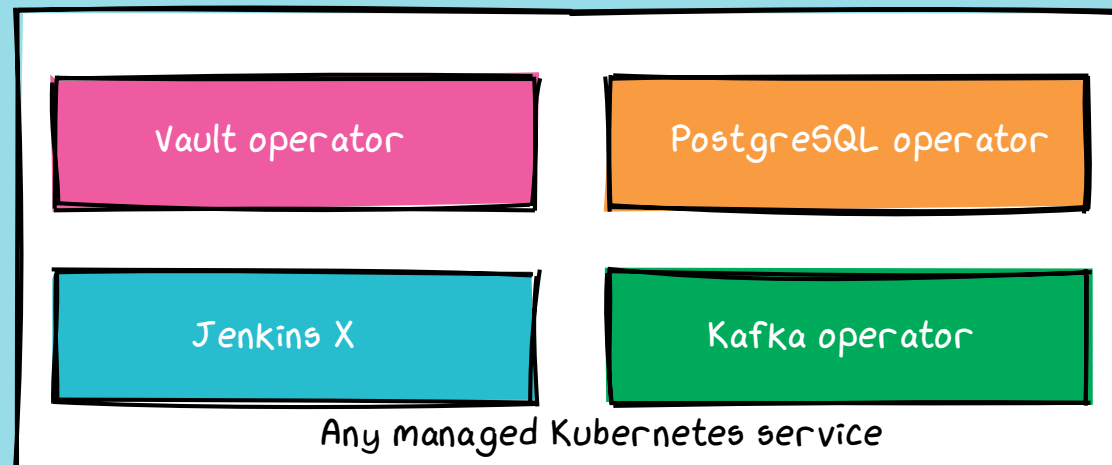
FALLING OUT OVER STATE MANAGEMENT



Cloud vendors perfected the use of proprietary services based on open source products, putting companies creating those products on notice.

In 2016 MongoDB Inc. responded with the Server Side Public License.

A PORTABLE STACK



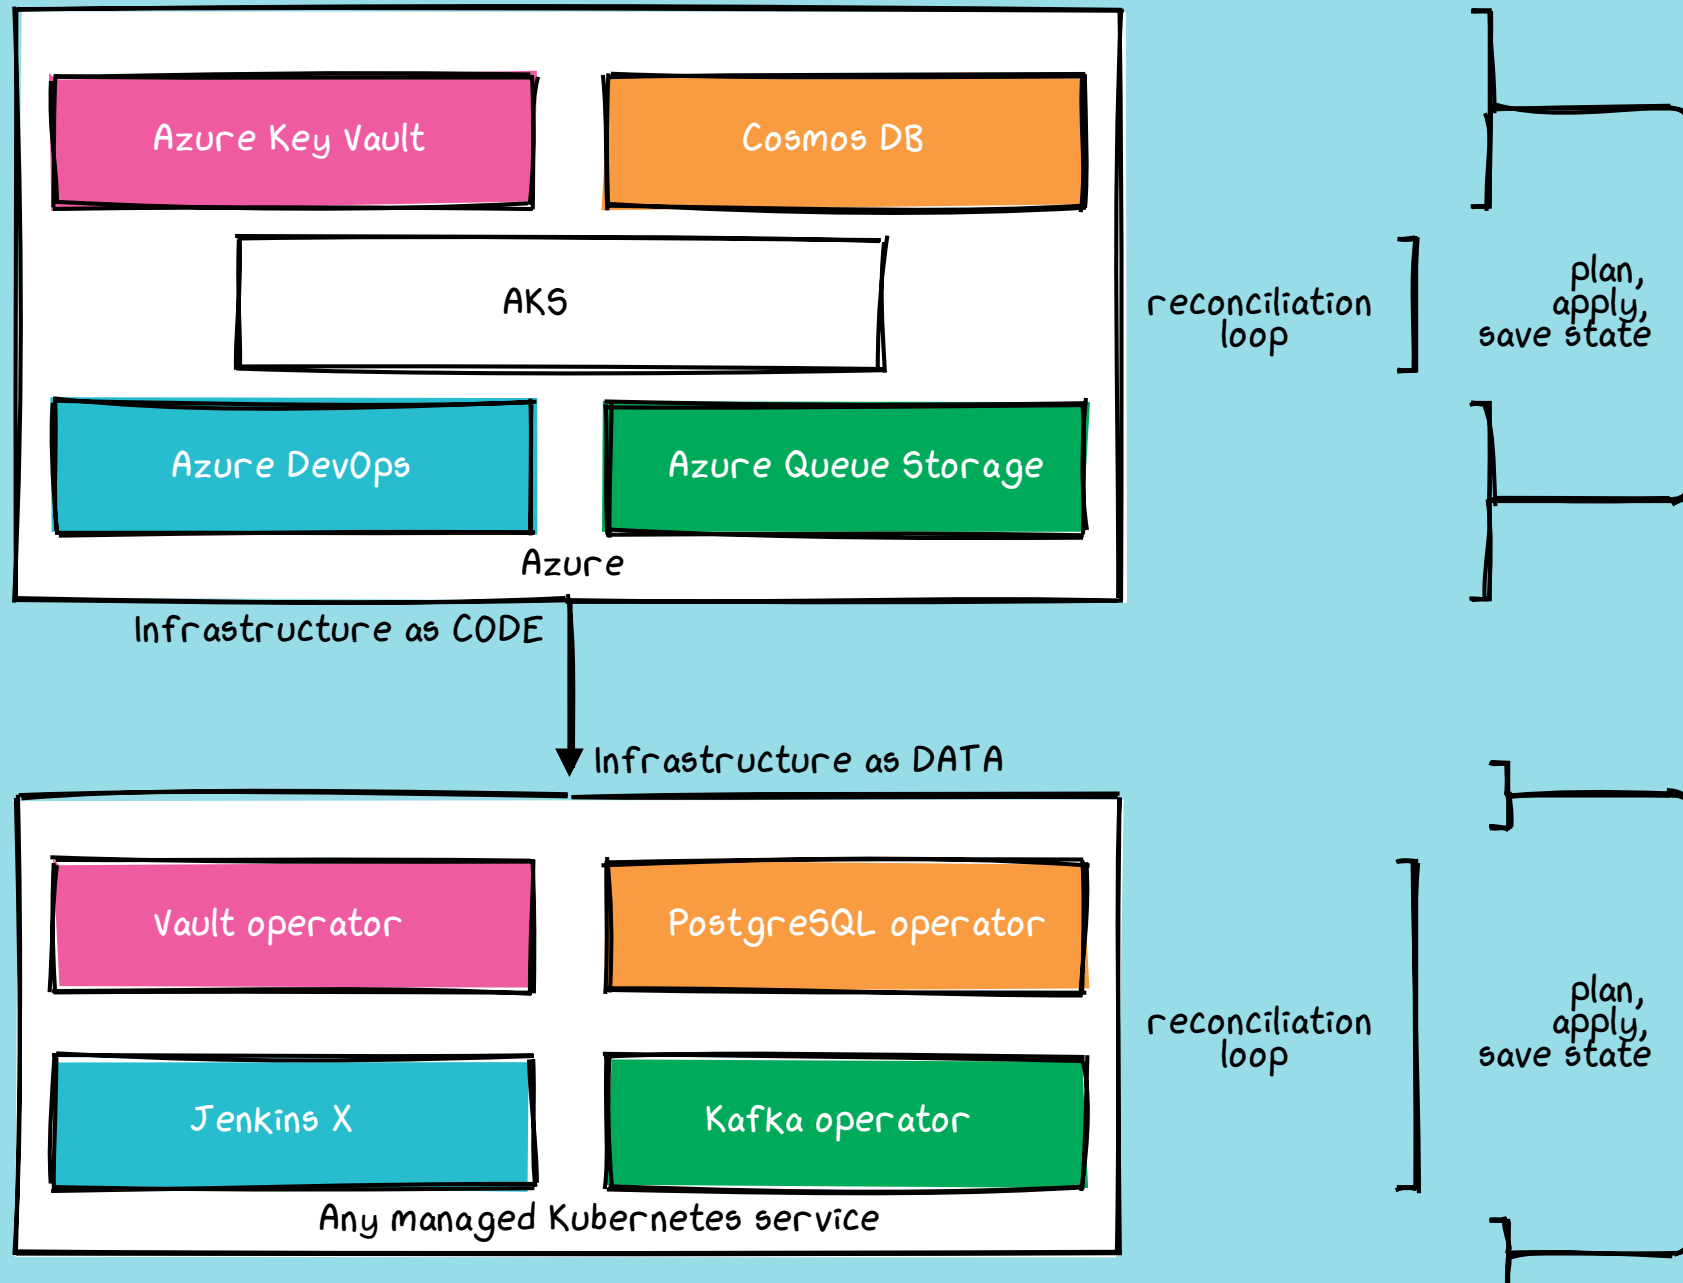
INFRASTRUCTURE AS ~~CODE~~ DATA

Declarative configuration is about treating infrastructure as data, which is more portable than code, and enables workflows that manipulate desired state based on policy.

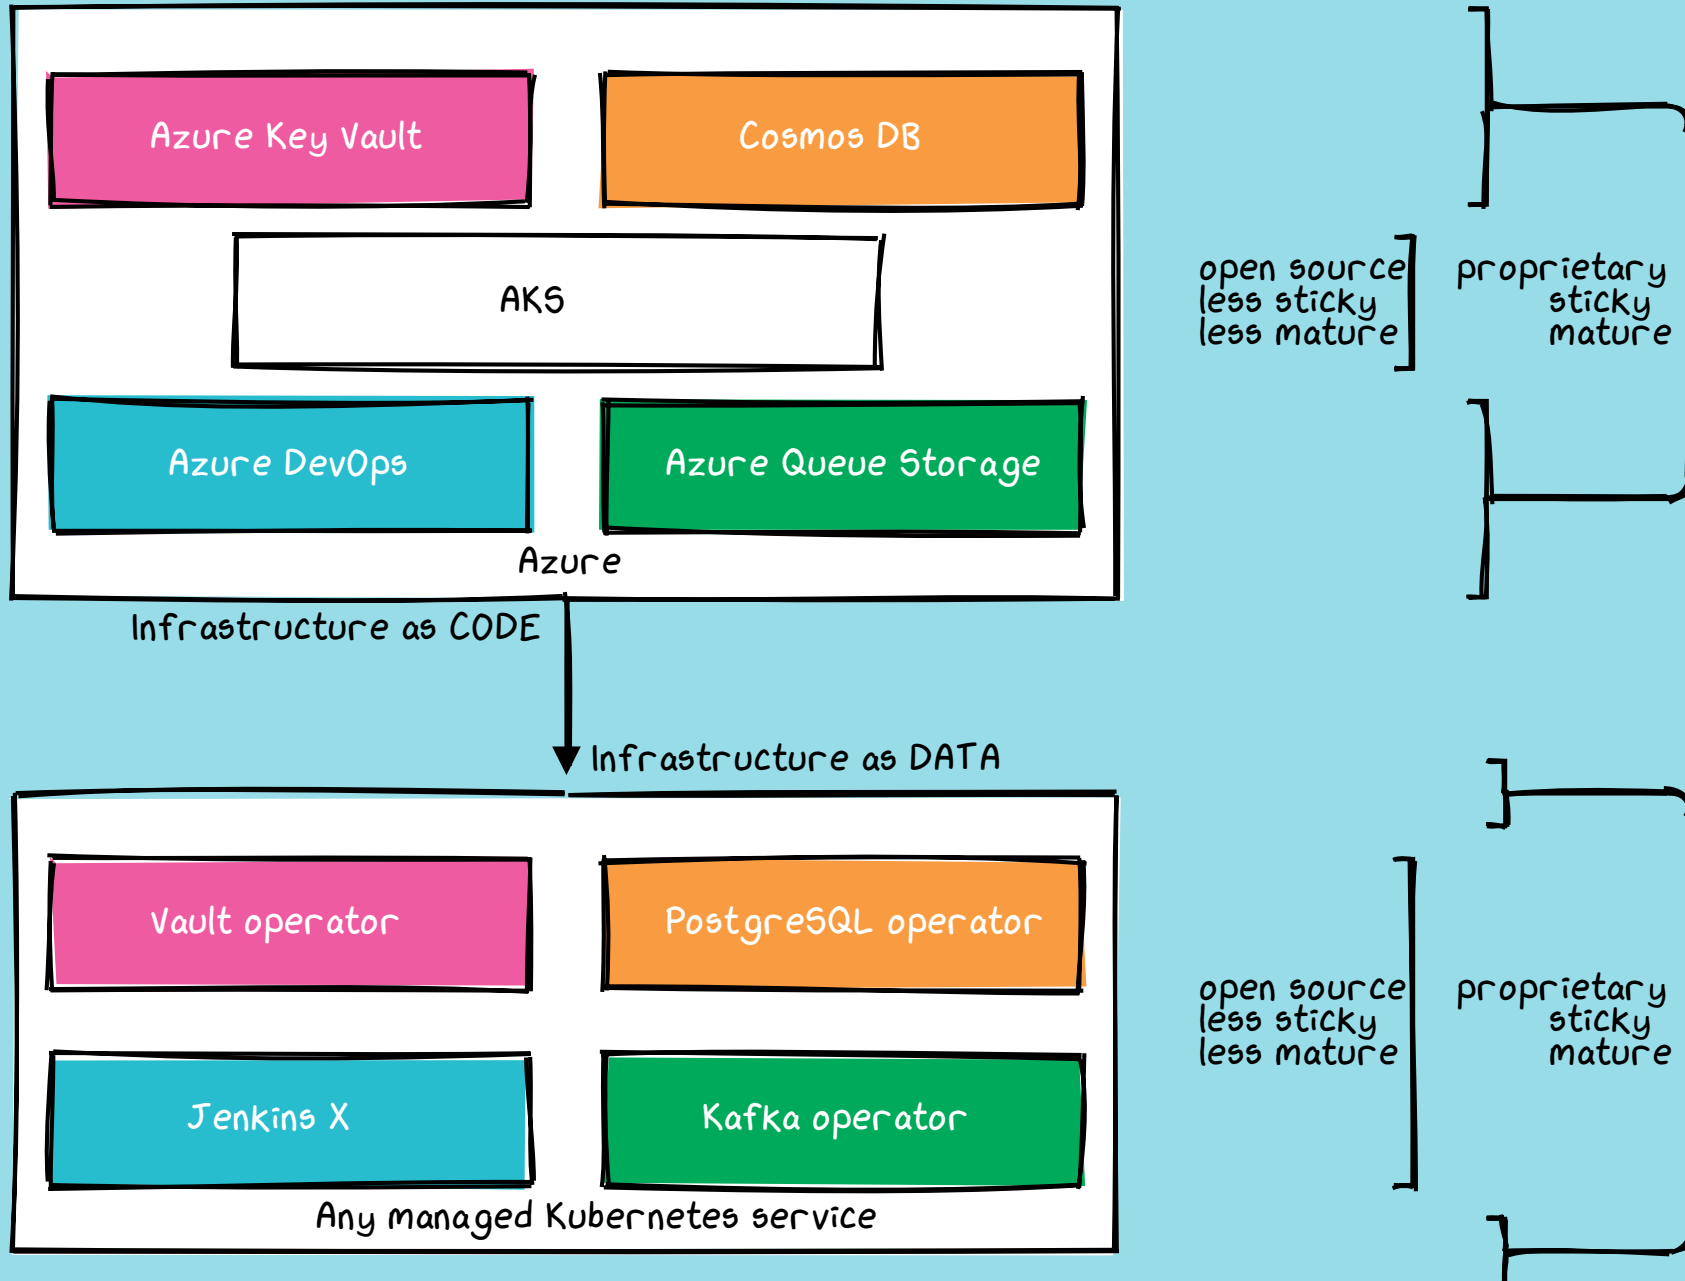
– Kelsey Hightower (2019)

Source: @kelseyhightower on Twitter, 21 August 2019.

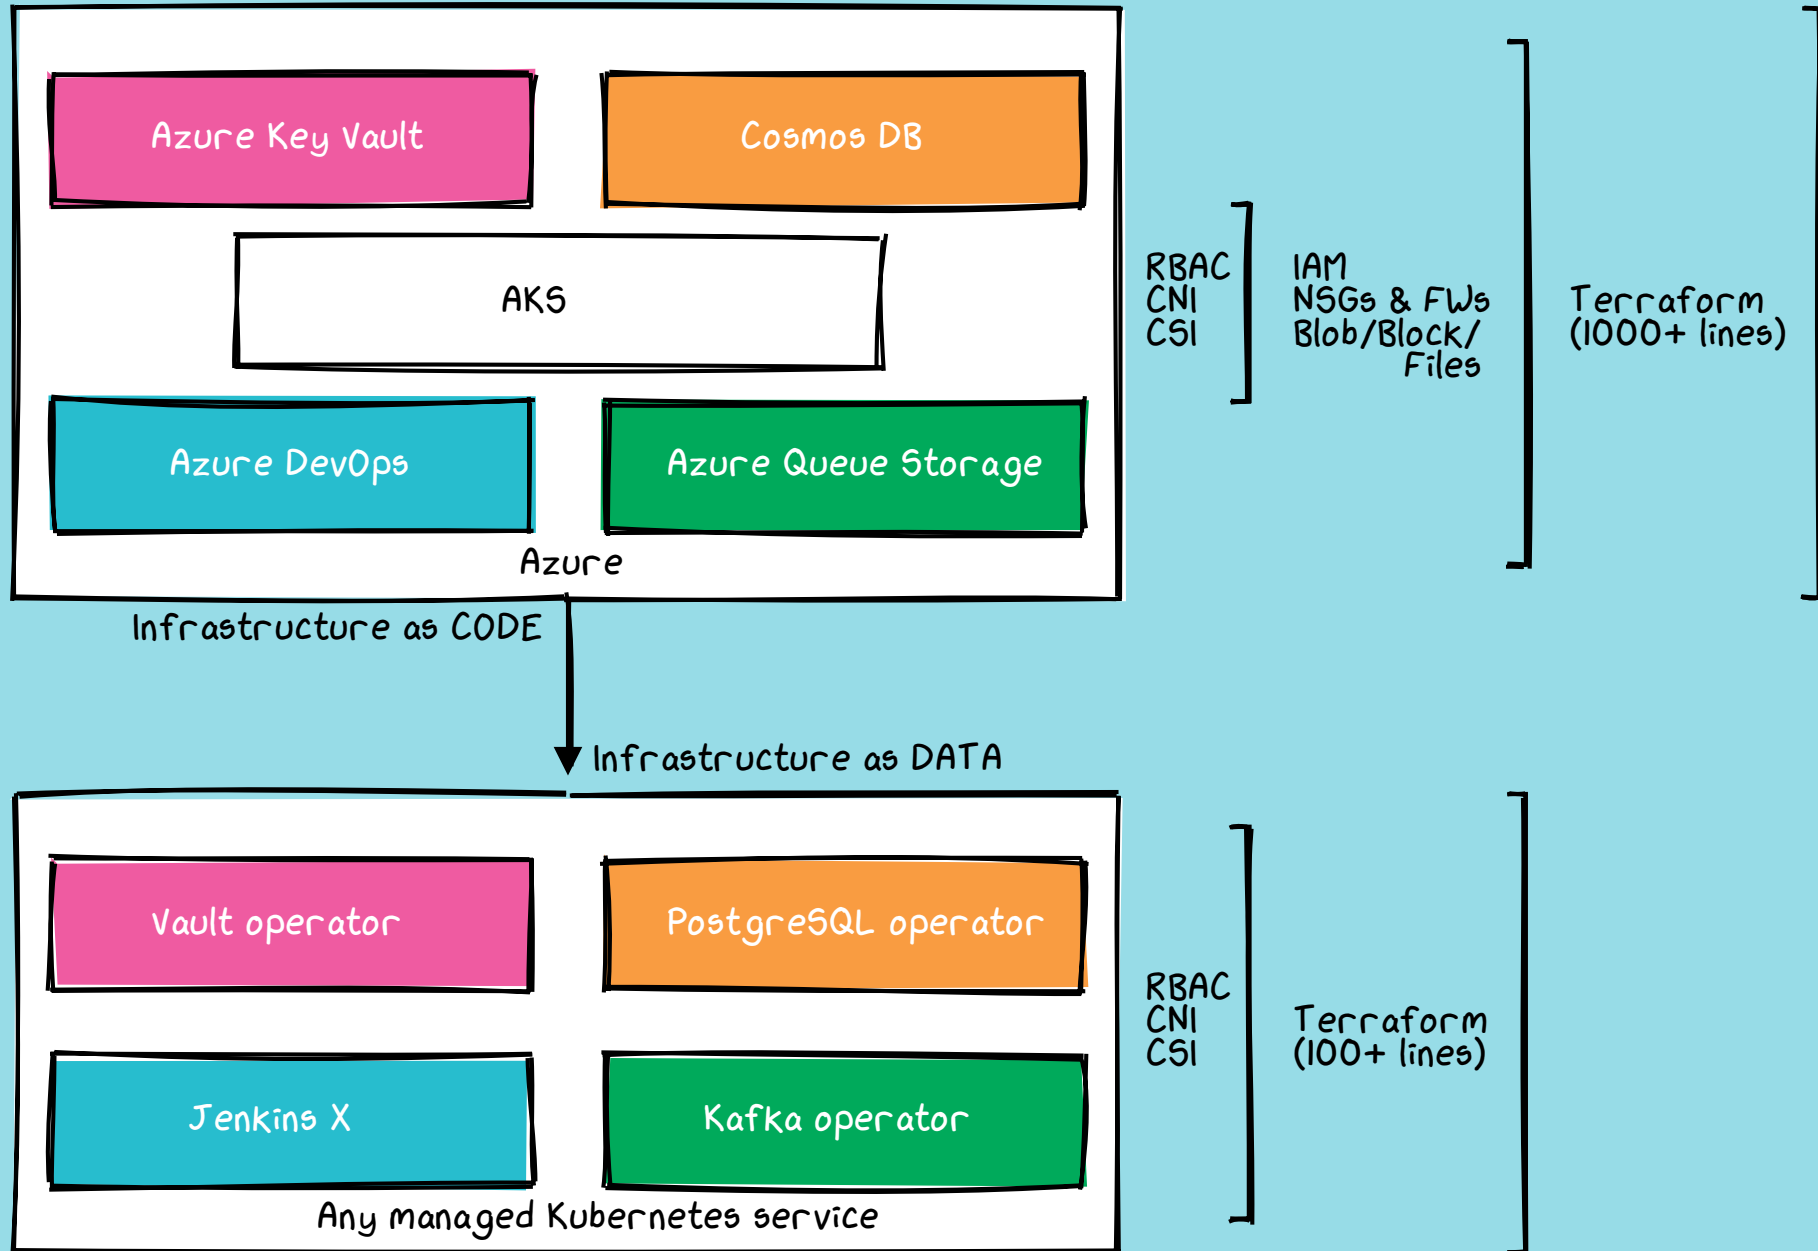
DESIRED STATE



TRADE-OFFS



CONFIGURATION



NATIVE RESOURCES

Can I trust custom resources to create and manage the lifecycle of objects native to the platform?

Yes, this is something resources like `PersistentVolumeClaim` and `Service` have done for a long time, dynamically creating storage volumes and load balancers respectively.

SERVICE LEVEL

Can the PostgreSQL operator match the availability and durability guarantees of managed SQL Server DBs?

Not today, no. It is worth considering, though, that:

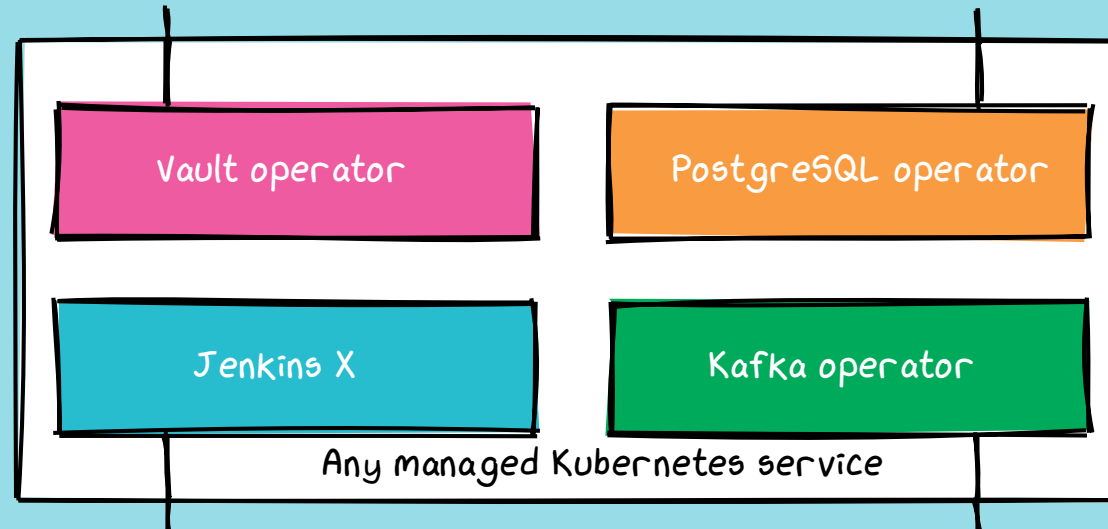
- * RDS and the operator use the same Azure storage primitives
- * Until this improves, there's Microsoft's Azure Service Operator for Kubernetes (EventHub, Azure SQL, CosmosDB, Storage Accounts)

Source: cloudblogs.microsoft.com/opensource/2020/06/25/announcing-azure-service-operator-kubernetes/

A PLAUSIBLE FUTURE, BUT NOT A DONE DEAL

The CoreOS Vault operator depends on their etcd operator

Two choices: Crunchy and Zalando's Spilo
MySQL has a first-party operator by Oracle
There is also CNCF graduate project Vitess



Note that this is not "Hudson"
UI requires VMware Octant with
octant-jx plugin

First-party operator by Confluent
One promising alternative is the
RabbitMQ Cluster Operator

WHAT'S IN IT FOR US?

As with multicloud, the value proposition for the consumer can seem lacklustre, but there are tangible benefits

Create one path to production for all workloads (not three separate ones for containers, functions and VMs)

Enabling zero trust and mutual TLS is much easier, and again you only do it once

Continuous reconciliation based on policies you define beats setting desired state at the start and hoping for the best

SUMMARY

Persistent cloud storage is just as reliable when claimed by a pod

The core Kubernetes interfaces will continue to mature because everyone in the industry has a stake in it

The operator pattern is set to thrive because software companies and second-tier cloud vendors depend on it

Of all the large, distributed systems found in modern application architectures, only Kubernetes promises a significant reduction in complexity in return

THANK YOU

Slides built with Markdeck

GitHub [gerald1248/stateful-kubernetes-slides](https://github.com/gerald1248/stateful-kubernetes-slides)

Twitter @03spirit

ThoughtWorks®